

Labwork_4

JACOB Mathieu

06/10/2025

1 Improve to use 2D blocks

We take the labwork 3 code and we modify it to use 2D blocks. We can find the modification in the course presentation. These modifications are that we use 2 variables which are x and y for the height and width. With this, we don't need to reshape the image before using it. We also launch the kernel in 2D that's why the functions for grayscale take vectors as parameters for blocks and thread per blocks.

```
#### GPU Blocks & Graph ####

from PIL import Image
import time
import numpy as np
from numba import cuda, config
import matplotlib.pyplot as plt
config.CUDA_ENABLE_PXXVITLINK = 1

# Function to convert in a grey image
@cuda.jit # for GPU
def greyscale_gpu(src, dst):
    tid1, tid2 = cuda.grid(2)
    if tid1 < src.shape[0] and tid2 < src.shape[1]: # Image limit
        r = src[tid1, tid2, 0]
        g = src[tid1, tid2, 1]
        b = src[tid1, tid2, 2]
        gy = np.uint8((r + g + b) / 3) # Average of pixels
        dst[tid1, tid2] = gy # The 3 pixels become grey

img = Image.open("../content/image.PNG").convert("RGB") # Load image
rgb_array = np.array(img, dtype=np.uint8)

dst = np.zeros((rgb_array.shape[0], rgb_array.shape[1]), dtype=np.uint8) # Allow memory
block_sizes = [1, 2, 4, 8, 12, 16, 20, 24, 28, 32] # Different block size values
times = [] # Timer

d_src = cuda.to_device(rgb_array) # Input
d_dst = cuda.device_array_like(dst) # Output

# The Loop
for threads_per_block in block_sizes:

    blocks_x = int(np.ceil(rgb_array.shape[0] / threads_per_block))
    blocks_y = int(np.ceil(rgb_array.shape[1] / threads_per_block))

    start_time = time.time() # Timer
    greyscale_gpu[(blocks_x, blocks_y), (threads_per_block, threads_per_block)](d_src, d_dst) # Function for greyscaling
    end_time = time.time()

    d_dst.copy_to_host(dst)

    time_ = end_time - start_time # Timer
    times.append(time_)
    print(f"Block size = {threads_per_block} x {threads_per_block} -> Time taken = {time_:.6f} sec")

# Graph
plt.figure(figsize=(7,4))
plt.plot(block_sizes, times, marker='o', linestyle='-', linewidth=2)
plt.xlabel('Block size n x n')
plt.ylabel('Execution time (sec)')
plt.title('Block size vs Time')
plt.grid(True)
plt.show()
```

Figure 1: Code for grayscale with GPU in 2D blocks

```

Block size = 1 x 1 -> Time taken = 0.077452 sec
Block size = 2 x 2 -> Time taken = 0.000418 sec
Block size = 4 x 4 -> Time taken = 0.000101 sec
Block size = 8 x 8 -> Time taken = 0.000087 sec
Block size = 12 x 12 -> Time taken = 0.000084 sec
Block size = 16 x 16 -> Time taken = 0.000075 sec
Block size = 20 x 20 -> Time taken = 0.000071 sec
Block size = 24 x 24 -> Time taken = 0.000087 sec
Block size = 28 x 28 -> Time taken = 0.000072 sec
Block size = 32 x 32 -> Time taken = 0.000069 sec

```

Figure 2: Result time taken with different block size values




Figure 3: Graph block size 2D vs time

We can observe on the graph the same phenomenon as the previous labwork which is a decrease in execution time with the increase in block size.

2 Comparison 2D and 1D

Here is the graph of time taken for different block size values in 1D :




Figure 4: Graph block size vs time

Thus, we can see a decrease in execution time with the increase in block size in the two graphs. However and except for the first value, we observe that the time taken for 2D blocks is significantly lower than time taken for 1D block. This is best seen if we compare the result instead of the graph :

```

Block size = 1 -> Time taken = 0.072486 sec
Block size = 4 -> Time taken = 0.000708 sec
Block size = 16 -> Time taken = 0.000238 sec
Block size = 32 -> Time taken = 0.000137 sec
Block size = 64 -> Time taken = 0.000126 sec
Block size = 128 -> Time taken = 0.000123 sec
Block size = 256 -> Time taken = 0.000134 sec
Block size = 512 -> Time taken = 0.000134 sec
Block size = 1024 -> Time taken = 0.000130 sec

```

Figure 5: Result time taken with different block size values in 1D

```

Block size = 1 x 1 -> Time taken = 0.077452 sec
Block size = 2 x 2 -> Time taken = 0.000418 sec
Block size = 4 x 4 -> Time taken = 0.000101 sec
Block size = 8 x 8 -> Time taken = 0.000087 sec
Block size = 12 x 12 -> Time taken = 0.000084 sec
Block size = 16 x 16 -> Time taken = 0.000075 sec
Block size = 20 x 20 -> Time taken = 0.000071 sec
Block size = 24 x 24 -> Time taken = 0.000087 sec
Block size = 28 x 28 -> Time taken = 0.000072 sec
Block size = 32 x 32 -> Time taken = 0.000069 sec

```

Figure 6: Result time taken with different block size values in 2D